

PROBLEM STATEMENT:- TO PREDICT THE RAINFALL BASED ON VARIOUS FEATURES OF THE DATASET

```
In [48]: import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [49]: df=pd.read_csv(r"C:\Users\sneha\OneDrive\Desktop\Rainfall in India.csv")
df
```

Out[49]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7
...
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4

4116 rows × 19 columns



DATA PREPROCESSING

In [50]:

df.head()

Out[50]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4

In [51]:

df.tail()

Out[51]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	1
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	2

```
In [52]: df.isnull().any()
```

```
Out[52]: SUBDIVISION    False
YEAR                False
JAN                  True
FEB                  True
MAR                  True
APR                  True
MAY                  True
JUN                  True
JUL                  True
AUG                  True
SEP                  True
OCT                  True
NOV                  True
DEC                  True
ANNUAL               True
Jan-Feb             True
Mar-May             True
Jun-Sep             True
Oct-Dec             True
dtype: bool
```

```
In [53]: df.fillna(method='ffill',inplace=True)
```

```
In [54]: df.isnull().sum()
```

```
Out[54]: SUBDIVISION    0
YEAR                0
JAN                  0
FEB                  0
MAR                  0
APR                  0
MAY                  0
JUN                  0
JUL                  0
AUG                  0
SEP                  0
OCT                  0
NOV                  0
DEC                  0
ANNUAL               0
Jan-Feb             0
Mar-May             0
Jun-Sep             0
Oct-Dec             0
dtype: int64
```

In [55]: `df.describe()`

Out[55]:

	YEAR	JAN	FEB	MAR	APR	MAY	JU
count	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000
mean	1958.218659	18.957240	21.823251	27.415379	43.160641	85.788994	230.56797
std	33.140898	33.576192	35.922602	47.045473	67.816588	123.220150	234.89605
min	1901.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.400000
25%	1930.000000	0.600000	0.600000	1.000000	3.000000	8.600000	70.475000
50%	1958.000000	6.000000	6.700000	7.900000	15.700000	36.700000	138.900000
75%	1987.000000	22.200000	26.800000	31.400000	50.125000	97.400000	306.150000
max	2015.000000	583.700000	403.500000	605.600000	595.100000	1168.600000	1609.900000

In [56]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SUBDIVISION     4116 non-null   object
1   YEAR            4116 non-null   int64
2   JAN             4116 non-null   float64
3   FEB             4116 non-null   float64
4   MAR             4116 non-null   float64
5   APR             4116 non-null   float64
6   MAY             4116 non-null   float64
7   JUN             4116 non-null   float64
8   JUL             4116 non-null   float64
9   AUG             4116 non-null   float64
10  SEP             4116 non-null   float64
11  OCT             4116 non-null   float64
12  NOV             4116 non-null   float64
13  DEC             4116 non-null   float64
14  ANNUAL          4116 non-null   float64
15  Jan-Feb        4116 non-null   float64
16  Mar-May        4116 non-null   float64
17  Jun-Sep        4116 non-null   float64
18  Oct-Dec        4116 non-null   float64
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

In [57]: `df.columns`

Out[57]: Index(['SUBDIVISION', 'YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL', 'AUG', 'SEP', 'OCT', 'NOV', 'DEC', 'ANNUAL', 'Jan-Feb', 'Mar-May', 'Jun-Sep', 'Oct-Dec'], dtype='object')

```
In [58]: df.shape
```

```
Out[58]: (4116, 19)
```

```
In [59]: df['ANNUAL'].value_counts()
```

```
Out[59]: ANNUAL
790.5      4
770.3      4
1836.2     4
1024.6     4
1926.5     3
..
443.9      1
689.0      1
605.2      1
509.7      1
1642.9     1
Name: count, Length: 3712, dtype: int64
```

```
In [60]: df['Jan-Feb'].value_counts()
```

```
Out[60]: Jan-Feb
0.0      238
0.1       80
0.2       52
0.3       38
0.4       32
...
23.3      1
95.2      1
76.9      1
66.5      1
69.3      1
Name: count, Length: 1220, dtype: int64
```

```
In [61]: df['Mar-May'].value_counts()
```

```
Out[61]: Mar-May
0.0       29
0.1       13
0.3       11
8.3       11
11.5      10
..
246.3     1
248.1     1
151.3     1
249.5     1
223.9     1
Name: count, Length: 2262, dtype: int64
```

```
In [62]: df['Jun-Sep'].value_counts()
```


```
Out[62]: Jun-Sep
434.3      4
334.8      4
573.8      4
613.3      4
1082.3     3
..
301.6      1
380.9      1
409.3      1
229.4      1
958.5      1
Name: count, Length: 3683, dtype: int64
```

```
In [63]: df['Oct-Dec'].value_counts()
```

```
Out[63]: Oct-Dec
0.0        16
0.1        15
0.5        13
0.6        12
0.7        11
..
191.5      1
124.5      1
139.1      1
41.5       1
555.4      1
Name: count, Length: 2389, dtype: int64
```

EXPLORATORY DATA ANALYSIS

```
In [64]: df=df[['JAN', 'FEB', 'MAR', 'APR', 'DEC']]
sns.heatmap(df.corr(),annot=True)
plt.show()
```



```
In [65]: df.columns
```

```
Out[65]: Index(['JAN', 'FEB', 'MAR', 'APR', 'DEC'], dtype='object')
```

```
In [66]: x=df[["FEB"]]
y=df["JAN"]
```

Linear Regression

```
In [67]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33
                                                ,random_state=42)
```

```
In [68]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(X_train,y_train)
print(reg.intercept_)
coeff_=pd.DataFrame(reg.coef_,x.columns,columns=['coefficient'])
coeff_
```

9.650666612303553

Out[68]:

	coefficient
FEB	0.442278


```
In [69]: score=reg.score(X_test,y_test)
print(score)
```

0.1793580786264921


```
In [70]: predictions=reg.predict(X_test)
```

```
In [71]: plt.scatter(y_test,predictions)
```

Out[71]: <matplotlib.collections.PathCollection at 0x14f51966c50>





```
In [72]: df500=df[:][:500]  
sns.lmplot(x="FEB",y="JAN",order=2,ci=None,data=df500)  
plt.show()
```



```
In [73]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33)  
reg.fit(X_train,y_train)  
reg.fit(X_test,y_test)
```

```
Out[73]: ▾ LinearRegression  
LinearRegression()
```

```
In [74]: y_pred=reg.predict(X_test)
plt.scatter(X_test,y_test,color='black')
plt.plot(X_test,y_pred,color='red')
plt.show()
```



```
In [75]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 Score:",r2)
```

R2 Score: 0.14759039309699007

Ridge Model

```
In [76]: from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

```
In [77]: features= df.columns[0:5]
target= df.columns[-5]
```

```
In [78]: x=np.array(df['JAN']).reshape(-1,1)
y=np.array(df['FEB']).reshape(-1,2)
```

```
In [79]: x= df[features].values
y= df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3
                                                ,random_state=17)
```

```
In [80]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge=ridgeReg.score(x_train,y_train)
test_score_ridge=ridgeReg.score(x_test,y_test)
```


```
In [81]: print("\n Ridge Model:\n")
print("The train score for ridge model is{}".format(train_score_ridge))
print("The test score for ridge model is{}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is0.9999999999874192
The test score for ridge model is0.99999999998833

```
In [82]: lr=LinearRegression()
```

```
In [83]: plt.figure(figsize= (10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",
        ,markersize=5,color="red",label=r'Ridge;$\alpha=10$')
plt.plot(features,alpha=0.4,linestyle='none',marker='o',
        markersize=7,color="green",label='LinearRegression')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



Lasso Model

```
In [84]: print("\n Lasso Model:\n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is{}".format(test_score_ls))
```


Lasso Model:

The train score for ls model is 0.9999207747038827

The test score for ls model is0.9999206791315255

```
In [85]: pd.Series(lasso.coef_,features).sort_values(ascending=True).plot(kind="bar")
```

Out[85]: <Axes: >




```
In [86]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,1,10],
                  random_state=0).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

0.9999999999999921

0.9999999999999921

```
In [87]: plt.figure(figsize= (10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",
        ,markersize=5,color='red',label=r'Ridge;$\alpha=10$')
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',
        ,markersize=6,color='blue',label=r'lasso;$\alpha=10$')
plt.plot(features,alpha=0.4,linestyle='none',marker='o',
        markersize=7,color="green",label='LinearRegression')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



ElasticNet

```
In [88]: from sklearn.linear_model import ElasticNet
elnet=ElasticNet()
elnet.fit(x,y)
print(elnet.coef_)
print(elnet.intercept_)
print(elnet.score(x,y))
```

```
[9.99098574e-01 0.00000000e+00 3.02728910e-05 0.00000000e+00
 0.00000000e+00]
0.016258606966612632
0.9999992160905338
```

```
In [90]: y_pred_elastic = elnet.predict(x_train)
mean_squared_error=np.mean((y_pred_elastic - y_train)**2)
print(mean_squared_error)
```

```
0.0008816302333951303
```

CONCLUSION:-

THE SCORE OF LINEAR REGRESSION IS :- 0.1793580786264921

THE SCORE OF RIDGE MODEL IS :- 0.99999999998833

THE SCORE OF LASSO MODEL IS :- 0.999999999999992

THE SCORE OF ELASTIC NET IS :- 0.9999992160905338

By observing the score of all the models we can prefer Lasso Model to this dataset

```
In [ ]:
```